INFORMATICS INSTITUTE OF TECHNOLOGY

In Collaboration with

UNIVERSITY OF WESTMINSTER

Interpretable Legal Document Classification and Verdict Prediction using Transformer-based NLP Models

A Project Proposal by
Lithila Ginige.

Supervised by
Mr. Shavin Fernando

September 2025

ABSTRACT

Problem:

The increasing reliance on Artificial Intelligence in the legal sector has created opportunities for automating judgment prediction based on legal case documents. However, existing Transformer-based NLP models—while highly accurate—operate as opaque "black boxes," providing limited insight into their decision-making process. This lack of interpretability undermines their practical adoption in judicial and advisory contexts, where transparency and reasoning are essential. Moreover, current models face challenges in handling long, complex legal documents and aligning machine-generated explanations with human legal reasoning.

Methodology:

This research proposes an interpretable legal document classification and verdict prediction framework using Transformer-based NLP models such as **BERT**, **RoBERTa**, and **Legal-BERT**. The system analyzes legal documents from the **European Court of Human Rights (ECHR)** dataset to predict case verdicts (violation or no violation). To ensure transparency, several **Explainable AI (XAI)** techniques—**SHAP**, **LIME**, and **Captum Integrated Gradients**—are applied to generate human-readable explanations highlighting influential text spans. Comparative experiments are conducted to assess performance, interpretability, and fairness across models. A **Streamlit-based web interface** is developed to visualize verdict predictions and explanations interactively.

Initial Results and Expected Contributions:

Preliminary experiments demonstrate that domain-specific models such as Legal-BERT outperform general models in accuracy, while SHAP-based explanations show greater faithfulness and interpretability compared to attention-based visualizations. The expected outcome is a transparent and trustworthy AI system that not only predicts verdicts accurately but also justifies its reasoning in legally meaningful terms. The study contributes to both the **technical domain** (multi-model XAI evaluation) and the **research domain** (human-aligned interpretability metrics for legal NLP).

Subject Descriptors :

- Computing methodologies → Natural language processing → Neural networks → Transformer models
- Information systems → Information retrieval → Text mining → Legal document analysis

- Computing methodologies → Machine learning → Explainable artificial intelligence

Keywords:

Legal judgment prediction, Transformer models, Explainable AI, Legal-BERT, SHAP, Interpretability

TABLE OF CONTENTS

CHAPTER 01: INTRODUCTION

1.1 Chapter Overview

This chapter provides an introduction to the research project titled "Interpretable Legal Document Classification and Verdict Prediction using Transformer-based NLP Models." It presents the problem background and the challenges in automating legal judgment prediction, followed by a detailed problem definition and motivation for the research. The chapter also discusses related existing work, identifies the research gap, outlines the proposed contributions, and states the aim, objectives, scope, and challenges of the project. Finally, it summarizes the hardware and software requirements that will support the development and experimentation phases.

1.2 Problem Background

The legal domain generates an immense amount of textual data daily, including court judgments, case descriptions, and legal arguments. Extracting meaningful insights and predicting outcomes from such unstructured text has become a growing challenge for both legal practitioners and researchers. The complexity of legal language, characterized by lengthy documents, domain-specific terminology, and intricate reasoning, makes automated processing particularly demanding.

In recent years, Natural Language Processing (NLP) has shown significant progress through the introduction of Transformer-based architectures such as BERT, RoBERTa, and Legal-BERT, which have outperformed traditional machine learning models in various text classification tasks. These models have also been applied in the legal domain to perform tasks such as judgment prediction, article citation, and legal summarization.

However, while these models achieve high predictive accuracy, they often operate as "black boxes," making it difficult for legal professionals to understand why a particular decision was predicted. This lack of transparency raises ethical and legal concerns regarding the accountability and trustworthiness of AI-assisted judicial systems. Therefore, developing a model that not only predicts legal verdicts but also provides clear, interpretable explanations is a crucial step toward building trust in legal AI systems.

1.3 Problem Definition

Although existing Transformer-based models can predict case outcomes with high accuracy, their predictions are typically opaque, offering little to no insight into the reasoning behind them. Current systems also struggle to handle lengthy legal documents effectively and fail to align machine-generated explanations with human legal reasoning. Consequently, there is a need for an interpretable, reliable, and fair legal judgment prediction framework capable of justifying its predictions through explainable AI (XAI) techniques.

1.3.1 Problem Statement

Existing Transformer-based legal judgment prediction models lack transparency and fail to provide interpretable explanations aligned with human legal reasoning.

1.4 Research Motivation

Legal professionals require decision support tools that are both accurate and interpretable. While Transformer-based NLP models have achieved state-of-the-art performance, their lack of explainability prevents their adoption in sensitive domains such as law. Moreover, ensuring fairness, handling long legal documents, and providing human-aligned reasoning remain unsolved challenges. Addressing these gaps is technically demanding, as it requires integrating deep learning architectures with interpretable AI mechanisms such as SHAP, LIME, and attention visualization, while maintaining predictive accuracy. This research aims to bridge that gap by developing a transparent legal AI system that enhances understanding, fairness, and trust.

1.5 Existing Work

Citation	**Summary**	**Limitation**	**Contribution**
Turan & Küçüksille (2025)	Used Transformer + SHAP to analyze Turkish Constitutional Court decisions.	Explanations are not always legally intuitive.	Demonstrated interpretability in legal decision prediction using XAI.
Ramani & Venugopalan (2025)	Applied BERT with XAI to classify rhetorical roles in Indian legal judgments.	Lacks evaluation with legal experts.	Improved transparency in legal text segmentation.
Gong et al. (2024)	Reviewed Transformer-based LJP and highlighted limitations of attention-based explanations.	Attention mechanisms weakly explanatory.	Proposed hybrid symbolic-neural XAI as future direction.
Chalkidis et al. (2022)	Applied Legal-BERT for paragraph-level role prediction with gradient-based XAI.	Lacks human-in-the-loop evaluation.	Introduced domain-adapted Legal-BERT with interpretability focus.

Wang et al. (2023)	Developed Case-Aware Transformer with attention visualization for explainable LJP.	Struggles with long documents and truncation.	Improved case-specific contextualization for LJP.

1.6 Research Gap

Most existing research on legal judgment prediction focuses primarily on improving accuracy but overlooks explainability and fairness. Attention visualization is commonly used as a proxy for explanation, but it has been shown to be insufficient for understanding model reasoning. Furthermore, many models are unable to process lengthy legal documents effectively, leading to loss of context and reduced performance. Few studies have attempted a systematic comparison of multiple XAI methods to evaluate which most closely aligns with human legal interpretation. This project addresses these gaps by developing an interpretable legal document classification and verdict prediction framework that combines multiple Transformer architectures with state-of-the-art XAI techniques and evaluates their interpretability and fairness using the ECHR dataset.

1.7 Contribution to the Body of Knowledge

Problem Domain Contribution:

The project will contribute an interpretable and explainable system capable of predicting verdicts from legal case documents using Transformer-based NLP models. It will also demonstrate how XAI methods such as SHAP, LIME, and Captum can be applied to provide meaningful, human-understandable justifications for legal predictions.

Research Domain Contribution:

This research will contribute to the academic field by providing a comparative analysis of Transformer models (BERT, RoBERTa, Legal-BERT) and XAI techniques in the context of legal NLP. It will evaluate the alignment of machine-generated explanations with human legal reasoning, offering novel insights into model interpretability, fairness, and explainability within the legal domain.

1.8 Research Challenges

- Handling long and complex legal documents exceeding standard Transformer token limits.
- Integrating XAI techniques while maintaining model accuracy.
- Quantitatively evaluating the quality of explanations.
- Addressing potential bias or unfairness in model predictions.
- Ensuring scalability and efficient computation using limited hardware resources.

1.9 Research Questions

1. How can Transformer-based NLP models be used to accurately predict legal verdicts from court documents?
2. Which XAI techniques (SHAP, LIME, Captum, attention visualization) provide the most legally intuitive explanations for verdict prediction?
3. How does document segmentation (facts, arguments, decisions) affect model accuracy and interpretability?
4. How can fairness and bias be evaluated and mitigated in Transformer-based legal judgment prediction systems?

1.10 Research Aim

The aim of this project is to design, develop, and evaluate an interpretable legal document classification and verdict prediction system using Transformer-based NLP models and Explainable AI techniques.

1.11 Research Objectives

Objective Code	Research Objective	Description	Learning Outcomes Mapped	Research Question Mapped
R01	To conduct an in-depth literature review on Transformer-based legal judgment prediction and explainable AI techniques.	Understand the state of the art and identify research gaps.	LO1, LO4	RQ1
R02	To preprocess and segment the ECHR dataset for legal document classification.	Clean, tokenize, and segment documents into meaningful sections.	LO2	RQ1, RQ3
R03	To develop and fine-tune multiple Transformer models (BERT, RoBERTa, Legal-BERT) for verdict prediction.	Build and evaluate different architectures on the ECHR dataset.	LO1, LO3	RQ1
R04	To integrate Explainable AI techniques (SHAP, LIME, Captum,	Generate token-level and feature-level explanations for predictions.	LO2, LO3	RQ2

	attention visualization) for model interpretability.			
R05	To assess fairness, human alignment, and interpretability quality across models.	Evaluate models on both performance and explainability metrics.	LO4	RQ2, RQ4
R06	To design a web-based interface demonstrating model predictions and explanations.	Create a user-friendly demonstration for legal experts.	LO5	RQ1, RQ2

1.12 Project Scope

1.12.1 In Scope

- Use of the ECHR dataset for training and evaluation.
- Implementation of Transformer models (BERT, RoBERTa, Legal-BERT).
- Integration of XAI methods (SHAP, LIME, Captum, attention visualization).
- Evaluation based on interpretability, fairness, and performance metrics.
- Development of a web-based interface for showcasing predictions and explanations.

1.12.2 Out of Scope

- Multilingual or non-English legal documents.
- Real-time deployment within judicial environments.
- Integration with legal databases or external APIs.
- Cross-jurisdictional legal reasoning (focus limited to ECHR cases).

1.13 Hardware & Software Requirements

1.13.1 Hardware Requirements

Component	**Specification**
CPU	Intel Core i7 / AMD Ryzen 7
GPU	NVIDIA RTX 3060 / Colab Pro GPU (12 GB VRAM)

RAM	Minimum 16 GB
Storage	100 GB for dataset, models, and logs

1.13.2 Software Requirements

Tool / Framework	**Version / Description**
Python	3.10+
PyTorch	2.0+
HuggingFace Transformers	Latest
SHAP / LIME / Captum	XAI toolkits
Streamlit	Frontend web demo
Google Colab / Kaggle	Cloud-based training
Git / GitHub	Version control
LaTeX / Overleaf	Paper preparation

1.14 Chapter Summary

This chapter introduced the motivation and objectives behind building an interpretable legal document classification and verdict prediction system. It highlighted the current limitations in Transformer-based legal AI models, particularly their lack of interpretability, and presented how the proposed research aims to address these gaps using Explainable AI techniques. The research questions, challenges, and objectives were defined, setting the foundation for the literature review and methodology described in the following chapters.

CHAPTER 02: LITERATURE REVIEW

2.1 Chapter Overview

This chapter presents a comprehensive review of existing research relevant to legal judgment prediction (LJP) and explainable artificial intelligence (XAI) within the legal domain. It begins by describing the problem domain and importance of interpretability in legal NLP. Subsequently, it reviews existing work using Transformer-based models for legal document analysis, highlighting their strengths and limitations. The chapter also discusses the datasets commonly used for judgment prediction, evaluation methodologies, and key challenges in interpretability and fairness. Finally, it identifies the research gaps that form the foundation for the proposed study.

2.2 Problem Domain

Legal judgment prediction (LJP) refers to the automatic prediction of court verdicts or outcomes based on textual case documents. The problem lies at the intersection of **Natural Language Processing (NLP)** and **computational law**, where models must interpret complex legal narratives, extract relevant facts, and simulate reasoning patterns similar to human judges.

2.2.1 Overview of Legal Judgment Prediction

The task typically involves classifying a case into verdict categories (e.g., violation/no violation) or predicting associated legal articles. Unlike ordinary text classification, LJP demands understanding of structured legal reasoning, precedent-based logic, and hierarchical document structures. Early LJP models relied on traditional machine learning approaches such as Support Vector Machines (SVMs) and Logistic Regression. However, these methods were limited by feature engineering and an inability to capture deep contextual semantics.

The advent of **Transformer architectures** such as BERT, RoBERTa, and Legal-BERT revolutionized legal NLP by leveraging attention mechanisms and large-scale pretraining, enabling context-aware understanding of legal text.

2.2.2 Importance of Explainable AI in Legal Systems

Despite their accuracy, Transformer-based models suffer from a lack of interpretability. In legal contexts, every decision must be transparent and justifiable; hence, "black-box" AI models are ethically and legally problematic. **Explainable AI (XAI)** techniques such as SHAP, LIME, and attention visualization aim to provide insight into which textual elements influenced a model’s decision. These explanations are crucial for promoting **trust, fairness, and accountability** in AI-assisted legal systems.

2.2.3 Challenges in Legal Document Processing

Legal documents are typically long, unstructured, and contain domain-specific language. The **length limitation of Transformer models (usually 512 tokens)** makes it difficult to capture the entire context. Additionally, legal reasoning often depends on implicit relationships between facts, statutes, and prior judgments, which are difficult for models to represent. Moreover, there is no universally accepted quantitative measure for interpretability in legal AI, complicating evaluation of XAI techniques.

2.3 Existing Work

2.3.1 Literature Survey Table

Citation	Model / Approach	Dataset	XAI Technique	Contribution	Limitation	Link / Source
Turan & Küçüksille (2025)	Hybrid Transformer + SHAP	Turkish Constitutional Court Decisions	SHAP	Demonstrated interpretable verdict prediction using SHAP to visualize important text spans.	Explanations not legally intuitive or human-aligned.	IEEE Access
Ramani & Venugopalan (2025)	BERT with Legal Named Entity Masking	Indian Legal Judgments	XAI framework	Improved role classification transparency using masked entities.	No expert evaluation from legal professionals.	IEEE Proceedings
Gong et al. (2024)	Review of Transformer-based LJP	Multiple legal datasets	Attention analysis	Identified weaknesses in attention as explanation,	No empirical validation; theoretical analysis only.	Springer

				suggesting hybrid symbolic-neural XAI.		
Chalkidis et al. (2021)	Multilingual Legal-BERT	MultiEURLEX (23 languages)	Saliency mapping	Enabled multi-label classification for cross-lingual legal text.	Cross-domain interpretability inconsistent across languages.	ACL Findings
Zhong et al. (2020)	Hierarchical BERT + Attention	CAIL2018 (China)	Attention visualization	Introduced hierarchical model for large-scale LJP.	Attention not a true measure of explanation.	AAAI
Chalkidis et al. (2022)	Paragraph-level Legal-BERT	ECHR dataset	Gradient-based XAI (Integrated Gradients)	Demonstrated paragraph-level interpretability in legal judgments.	Lacks human-in-the-loop evaluation.	ACL
Wang et al. (2023)	Case-Aware Transformer	Chinese Legal Datasets	Attention-based visualization	Enhanced context sensitivity between facts and laws.	Fails on very long cases due to token truncation.	ScienceDirect
Zhong et al. (2022)	XAI evaluation metrics for LJP	CAIL2018	SHAP, Attention	Proposed metrics for legal intuitiveness of	Explanations still diverge from	COLING

				explanations.	lawyer reasoning.	

2.3.2 Critical Evaluation

Existing Transformer-based LJP systems have achieved high accuracy by capturing contextual semantics across lengthy legal texts. However, interpretability remains the key shortcoming. For example, **Zhong et al. (2020)** and **Wang et al. (2023)** employed attention mechanisms to visualize important tokens, but subsequent studies (Gong et al., 2024) demonstrated that attention alone is not equivalent to human-understandable reasoning. **Turan & Küçüksille (2025)** successfully applied SHAP to visualize token contributions, but their explanations lacked the legal interpretability required by practitioners.

Chalkidis et al. (2021, 2022) introduced domain-specific models such as Legal-BERT, improving contextual representation of legal terms. However, their explainability evaluation was limited, with few studies involving human experts. **Zhong et al. (2022)** emphasized this limitation by proposing quantitative metrics for measuring "legal intuitiveness," yet alignment between human reasoning and model attributions remains underexplored.

From a methodological perspective, most studies evaluate models on single datasets (e.g., CAIL or ECHR) without testing generalization across jurisdictions or fairness across demographic attributes. Moreover, while multiple XAI techniques exist, no comprehensive comparison has been made to determine which best supports legal interpretability. These limitations collectively motivate this project's multi-model, multi-XAI comparative framework.

2.4 Dataset Selection

This project utilizes the **European Court of Human Rights (ECHR) dataset**, one of the most widely used resources in legal NLP research. The dataset contains thousands of legal cases labeled according to articles of the European Convention on Human Rights, with binary verdicts (violation or no violation).

Key characteristics:

- **Language:** English
- **Documents:** Over 11,000 case texts
- **Structure:** Includes case facts, arguments, relevant articles, and final judgments
- **Labels:** Binary (violation / no violation)

- **Advantages:** High-quality annotation, transparency, and relevance for human rights law

Preprocessing plan:
Documents will be segmented into sections such as “Facts,” “Arguments,” and “Law.” Since many cases exceed the Transformer token limit, a **document chunking strategy** will be applied, where each chunk is encoded separately and aggregated using mean-pooling or attention-based fusion. This ensures better long-context handling without loss of key information.

2.5 Benchmarking and Evaluation

Benchmarking involves comparing multiple Transformer-based models and XAI methods on both predictive performance and interpretability.

2.5.1 Evaluation Metrics

- **Performance Metrics:** Accuracy, Precision, Recall, F1-score, AUC-ROC
- **Interpretability Metrics:** Faithfulness (how well explanations reflect model reasoning), Legal Intuitiveness (overlap with human rationales), and Explanation Stability (consistency across similar cases)
- **Fairness Metrics:** Group fairness (demographic parity) and case-based fairness (prediction consistency across similar fact patterns)

2.5.2 Comparative Analysis

Aspect	**BERT**	**RoBERTa**	**Legal-BERT**
Domain Adaptation	General domain	General domain, robust pretraining	Legal domain fine-tuned
Performance	Strong baseline	Higher accuracy with optimized pretraining	Best on legal corpora
Explainability	SHAP, LIME compatible	SHAP, Captum, attention visualization	SHAP, Captum, gradient-based methods
Limitation	Short-context	Requires large GPU resources	Domain-specific vocabulary bias

The project will integrate **SHAP**, **LIME**, and **Captum's Integrated Gradients** to visualize and compare token-level attributions. Each XAI method will be quantitatively assessed for **faithfulness** and **human-alignment** through overlap with legal rationales.

2.6 Chapter Summary

This chapter reviewed the state of research in Transformer-based legal judgment prediction and interpretability. It revealed that although models such as BERT, RoBERTa, and Legal-BERT have improved prediction accuracy, their decision-making processes remain opaque. Attention visualization, while common, provides limited interpretability and often fails to align with human legal reasoning. Additionally, challenges such as long-document handling, lack of fairness evaluation, and absence of standardized interpretability metrics persist.

These gaps justify the need for a comprehensive, interpretable legal verdict prediction system integrating multiple Transformer architectures with XAI methods. The next chapter outlines the research methodology adopted to design, develop, and evaluate this system.

CHAPTER 03: METHODOLOGY

3.1 Chapter Overview

This chapter presents the methodological framework adopted for the development of the proposed interpretable legal document classification and verdict prediction system. The methodology follows both a **quantitative experimental approach** for model development and evaluation, and a **qualitative approach** for analyzing interpretability and fairness aspects. It discusses the overall research philosophy, approach, methodological choices, data collection techniques, and system development process. The chapter also describes the project management plan, hardware/software resources, and risk mitigation strategies used to ensure successful completion of the research.

3.2 Research Methodology

The research follows the structure of **Saunders' Research Onion**, as shown in Table 3.1, to define and justify the philosophical, methodological, and strategic choices made throughout the study.

Table 3.1: Saunders' Research Onion

Layer	Choice (What is Being Used)	Justification (Why You Are Using It)
Research Philosophy	Pragmatism	Combines both quantitative and qualitative approaches, allowing for performance evaluation of models (quantitative) and interpretability analysis based on expert judgment (qualitative).
Approach to Theory Development	Deductive	The research tests hypotheses that Transformer-based models can effectively predict verdicts and that XAI techniques improve interpretability.
Methodological Choice	Mixed Methods	Employs quantitative metrics for model evaluation and qualitative

		insights for human interpretability and fairness assessment.
Research Strategy	Experimental	Models and XAI methods are tested under controlled experimental conditions using the ECHR dataset.
Time Horizon	Cross-sectional	Data from the ECHR dataset represents a static snapshot used for model training and evaluation.
Data Collection Method	Secondary Data (Public Dataset)	Uses pre-existing, publicly available legal datasets to ensure reproducibility and transparency.

3.3 Development Methodology

3.3.1 Requirement Elicitation Methodology

The requirements for this system were identified through:

- **Literature Review (Archival Research):** Analysis of previous studies on legal judgment prediction and explainable AI.
- **Observation:** Evaluation of how current legal AI models present outputs and explanations.
- **Self-Evaluation and Brainstorming:** Determination of system needs such as interpretability metrics, model visualization, and fairness testing.

This combination ensured that the requirements were both **research-oriented** and **practically feasible** within the FYP timeframe.

3.3.2 Design Methodology

The project follows an **Object-Oriented Analysis and Design Methodology (OOADM)**, focusing on modular and reusable code components for preprocessing, model training, explanation generation, and user interface visualization.

Key design components:

- **Data Preprocessing Module:** For cleaning and segmenting ECHR cases.
- **Model Module:** For fine-tuning Transformers (BERT, RoBERTa, Legal-BERT).
- **Explainability Module:** Integrates SHAP, LIME, Captum, and attention-based interpretability.
- **Evaluation Module:** For measuring accuracy, interpretability, and fairness.
- **UI Module:** Streamlit-based interface to visualize verdicts and explanations.

3.3.3 Programming Paradigm

The development will adopt **Object-Oriented Programming (OOP)** using **Python**, promoting modularity and maintainability. Each system component (data loader, model trainer, explainer, and evaluator) will be encapsulated as a separate class.

3.3.4 Testing Methodology

Testing is divided into **model-level** and **system-level** evaluations:

Testing Type	**Objective**	**Metrics / Evaluation Criteria**
Model Testing	Evaluate prediction accuracy of each Transformer model.	Accuracy, F1-score, AUC-ROC
XAI Testing	Assess the interpretability of SHAP, LIME, Captum explanations.	Faithfulness, Legal Intuitiveness
Fairness Testing	Evaluate bias in predictions across subgroups.	Disparate Impact Ratio
Prototype Testing	Ensure UI correctness and visualization functionality.	Usability, response time, visual clarity

3.3.5 Solution Methodology

The overall solution follows a structured **machine learning pipeline**, described below:

1. **Dataset Collection:**
    - Use the **ECHR dataset** (European Court of Human Rights).
    - Contain text documents labeled by verdict outcomes.

2. **Data Preprocessing:**
    - Clean HTML and noise, remove redundant sections.
    - Tokenize text using pre-trained tokenizer.
    - Segment documents into meaningful sections (facts, arguments).
    - Apply chunking for documents exceeding 512 tokens.
3. **Model Selection and Training:**
    - Fine-tune **BERT**, **RoBERTa**, and **Legal-BERT** on the preprocessed data.
    - Use stratified training and validation splits.
    - Apply early stopping to prevent overfitting.
4. **Explainability Integration:**
    - Generate token-level explanations using **SHAP**, **LIME**, and **Captum (Integrated Gradients)**.
    - Visualize attention maps and compare across models.
    - Quantify faithfulness and human-alignment.
5. **Evaluation:**
    - Compare models based on classification and interpretability metrics.
    - Perform error and fairness analysis.
6. **Web Interface Development:**
    - Build a **Streamlit** web application to visualize predictions, confidence scores, and token-level explanations interactively.
7. **Documentation and Publication:**
    - Summarize experiments and results for research paper publication.

3.4 Project Management Methodology

3.4.1 Gantt Chart

3.4.2 Schedule

Table 3.2: Project Timeline and Deliverables

Deliverable	Target Date
Project Proposal (Initial Draft)	26th September 2025
Project Proposal and Requirement Specification (Final Draft)	24th October 2025
Project Proposal and Requirement Specification (Final)	14th November 2025
Proof of Concept	14th November 2025
Design Document	2nd February 2026
Prototype and Interim Project Demo	2nd February 2026
Implementation and Testing	February – March 2026
Final Evaluation and Thesis Submission	1st April 2026
Minimum Viable Product	1st April 2026

3.5 Resource Requirements

3.5.1 Skills Required

- Deep learning (Transformer architectures)
- Explainable AI and interpretability methods
- Data preprocessing and NLP tokenization
- Streamlit web development
- Research writing and documentation

3.6 Risk Assessment and Mitigation

Risk	Severity (1–5)	Frequency (1–5)	Mitigation Strategy
Insufficient data or preprocessing errors	4	3	Use data augmentation and verify cleaning scripts.

Model overfitting	5	4	Apply regularization, early stopping, and cross-validation.
Hardware limitations	3	3	Use Google Colab or Kaggle GPUs for training.
Poor model interpretability	5	3	Evaluate multiple XAI techniques and tune parameters.
Ethical or fairness issues	4	2	Perform fairness audits and use bias detection tools.
Time management risks	3	3	Maintain Gantt schedule and iterative progress reviews.
Integration failures (UI)	3	2	Modular development and staged testing.

3.7 Chapter Summary

This chapter presented the methodological foundation of the research, structured according to Saunders' Research Onion. The study employs a **pragmatic mixed-methods approach** that integrates quantitative performance evaluation and qualitative interpretability analysis. The chapter outlined the experimental pipeline — from data preprocessing to model explainability and visualization — and described the project management strategy, resource allocation, and risk mitigation plan. The next stages will involve implementation, testing, and evaluation of the developed system in alignment with the defined objectives.